Space Collect MapReduce Processing Model

LeoDOS Project

February 25, 2026

This document describes the design of the Space Collect MapReduce Processing (SpaceCoMP) Model from an application programmer's perspective.

1 Overview

The MapReduce model has been widely adopted as a primitive for parallel computing of data-intensive workloads in compute clusters. Mappers process data in parallel, and Reducers consolidate, sort and aggregate the output from mappers. The MapReduce primitive is popular as it shields the application programmers from decisions and complexities around scheduling, orchestration, reliability, and data locality. It also works equally well with simple one-off batch jobs as well as complex streaming data processors in elaborate hierarchical DAG workflows.

The idea with SpaceCoMP is that in space, or more precisely in ISL-mesh-connected mega-constellations of satellites in LEO orbits, sensor data is generated in a distributed pattern as well. The data production may for instance be correlated with a projection of an area of interest (AOI) on Earth to observe something like a developing forest fire. Due to orbital dynamics, the satellites that cover this AOI changes constantly but predictably. When orchestrating sensor data collection and processing nodes, this dynamics has to be taken into account to avoid excessive data shuffling. As the satellites move in their orbits, the orchestration and scheduling decisions have to be re-evaluated. Communication in ISL mesh networks follow a hop-by-hop pattern best described by Manhattan distances. Thus distances between nodes both in terms of hops and physical distances of each link for each hop has to be considered. Routing decisions also need to consider that there are orbital cross-over events where some links are unreliable. Another fundamental difference to traditional MapReduce processing in ground data centers is that scheduling decisions need to be taken in a decentralized manner, as many concurrent jobs may be injected from any satellite in a constellation that is visible from a point on the ground. We refer to the satellite that a ground station communicates with as the line-of-sight (LOS) satellite.

The fundamental principles of SpaceCoMP are:

- Collectors are scheduled based on predicted current AOI projection (and load)

- Mappers are scheduled based on cost of transmitting and processing data from collectors

- Reducer is scheduled based on proximity to mappers (and load)

In addition to these constraints, the current load on each satellite is also considered to balance the load across many concurrent requests, potentially from many LOS entry points and with overlapping AOI.

Figure 1 depicts the high-level architecture of a SpaceCoMP workflow.

Application programmers may provide custom or use generic Collector, Mapper, and Reducer processors described below.

To simplify the discussion, we assume here that these components are implemented in Python and that the payloads being passed follow a JSON format (with the exception of binary data which is streamed separately).

2 Collector

The collectors implement a collect method with a JSON payload input parameter:

```python
def collect(payload):
```

Figure 1: SpaceCoMP Architecture

A payload `meta_data` field contains `meta_data` that travels with the job through all phases including:

data_id	Unique index across all collectors
data_size	Total number of collectors
job_start	Timestamp when job was submitted
jobid	Globally Unique Job Identifier
max_collect	How many collect tasks should be buffered before streaming to mapper
job_data	Application specific data

When a collector has collected a data record it is emitted with a yield statement, meaning it is ready to be processed by a mapper:

```
yield data
```

The data object may by any serializable valid JSON.

If binary data is collected, e.g. an image needs to be sent to mappers for processing the collector should return the following structure:

```
yield {"value": data, "_COMP_FILE_": {"name": fname, "stream": stream}}
```

Where **fname** is the filename and the stream is the binary file object stream. The binary stream can simply be an open file or be created from a bytearray with e.g:

```
stream = io.BytesIO(byte_arr)
```

The basic design is that the collector can focus on expensive IO processing and maybe some initial pre-processing of sensor data before streaming records to a mapper for more elaborate processing. That way the collection and pre-processing can be parallelized across data records even within the same collect/map pair.

3   Mapper

The mappers implement a `run_map` method with a JSON payload input parameter:

```
def run_map(payload):
```

This payload parameter has the same **meta_data** field as described above for the collector as well as:

data	The record the collector produced
files	Dictionary with keys being filenames of binary data set in the collector. The value is a stream that can be read from just like a local file.
end_collect	Boolean set to indicate whether this is the last record produced by the collector
collected_index	The ordering number of record collected
collector	The satellite that collected the data

Just like with the collector both basic JSON and binary streams can be emitted for reduce processing with a yield statement.

4 Reducer

The reducer implement a reduce method with a JSON payload array input parameter:

```
def reduce(payloads):
```

A payload is one map record and there may be many map records for each mapper. The last record from a mapper is marked with an **end_map** flag similar to the **end_collect** flag described above. The data received for each payload mimics the mapper structure:

data	The record the mapper produced
files	Dictionary with keys being filenames of binary data set in the collector. The value is a stream that can be read from just like a local file.
end_map	Boolean set to indicate whether this is the last record produced by the mapper
mapped_index	The ordering number of record mapped
mapper	The satellite that mapped the data

The reducer can produce JSON data or binary stream output the same way as the collectors and mappers but instead of yielding or streaming records it simply returns the result, e.g.:

```
return {"value": data, "_COMP_FILE_": {"name": fname, "stream": stream}}
```

Next, to illustrate how the processors above can be used to implement different applications we provide a "hello world" example (Word Count) and a couple of examples related to Earth Observation.

5 Combiner

The combiner implements a combine method with a data object of intermediate, collected data and a JSON payload array input parameter:

```
def combine(data,payloads):
```

The data parameter is an array of records yielded from the collector. The number of records depends on the total records yielded and the job meta data parameter *max_collect_records*.

To trigger the combine method to be called after a collect call is made, the collector should set the following instance variable.

```
def __init__(self):
    self.COMP_SKIP_MAP = True
```

This flag will cause the combiner to be called in the same satellite as the collector and the combined output will be sent directly to the reducer.

6   Interactions

The communication and data transfers between the LOS, Collector, Mapper and Reducer satellites are handled by the SpaceCoMP system. The actual transport protocol used is not defined here but is assumed to be some reliable and CCSDS compliant transport like DTN. Ground simulations may use HTTP or UDP. All messages are asynchronous as they may travel many hops via the ISL mesh to their destination satellite.

See Figure 2 for the key interactions between the SpaceCoMP components.

7   Word Count

To illustrate the simplest possible SpaceCoMP application we implemented the classical word count example. A mapper reads a chunk of a file and streams it line by line to the mapper that counts the words and streams the word counts to a reducer that produces the final word counts across all mappers.

See:

- `https://github.com/LeoDOS-Project/leopymr/blob/main/docker/collectors/doccollector.py`

- `https://github.com/LeoDOS-Project/leopymr/blob/main/docker/mappers/wordcountmapper.py`

- `https://github.com/LeoDOS-Project/leopymr/blob/main/docker/reducers/sumreducer.py`

8   SAR Example

Synthetic Aperture Radar (SAR) Images are popular in satellite observation applications as the pictures produced are less susceptible to atmospheric distortion and cloud obstruction.

The images are however large in size (up to 7GB for a single image) and some denoising is necessary.

In our example a custom mapper denoises SAR TIFF images using a deep (neural network) despeckling tool, and the denoised images are then sent to a custom mapper that does object detection with computer vision. The CV output (count of objects detected) is then sent to a simple reducer that just takes the counts from all mappers and aggregates them to produce the final results.

More info at: `https://github.com/LeoDOS-Project/leopymr/blob/main/usecases/sar/README.md`

9   MISR Example

Multi-Image(Frame) Super Resolution is a popular image processing algorithm used in Earth Observation use cases to combine many lower resolution images taken by satellites into a single higher resolution image.

In our example a custom collector reads a burst of images in dng format and converts them to png before streaming them to a custom mapper that implements the algorithm used in Android to create higher resolution images from a burst. The mappers then forwards the combined image to the reducer that takes all the images from the mapper to produce the final merged image output using the same algorithm.

More info at: `https://github.com/LeoDOS-Project/leopymr/blob/main/usecases/misr/README.md`

10   VJEPA Example

Video-Joint Embedding Predictive Architecture (V-JEPA) is a framework for self-supervised training with multi-modal data. The models are particularly useful in predicting future events from a stream of video frames. Unlike other object detection models it does not predict on a pixel level but rather via abstract semantic and latent parameter spaces making them more robust to new data as they learn *world models* or *physics* to make good guesses for situations previously not seen.

The VJEPA model can be run for a live stream of video frames where a series of frames are buffered and run through the predictor continuously to recognize events, in this example movement.

Since the results are produced already in the collector there is no need to run a mapper other than to concolidate the results before sending it on to the reducer in comnpressed form. This compression can be done locally, so this example makes use of the *SKIP_MAP* flag, and a combiner (the *mergecombiner*).

The VJEPA model we use is taken from huggingface so this use case also demonstrates how models can be pulled from huggingface and loaded onto the satellite runtime during deployment.
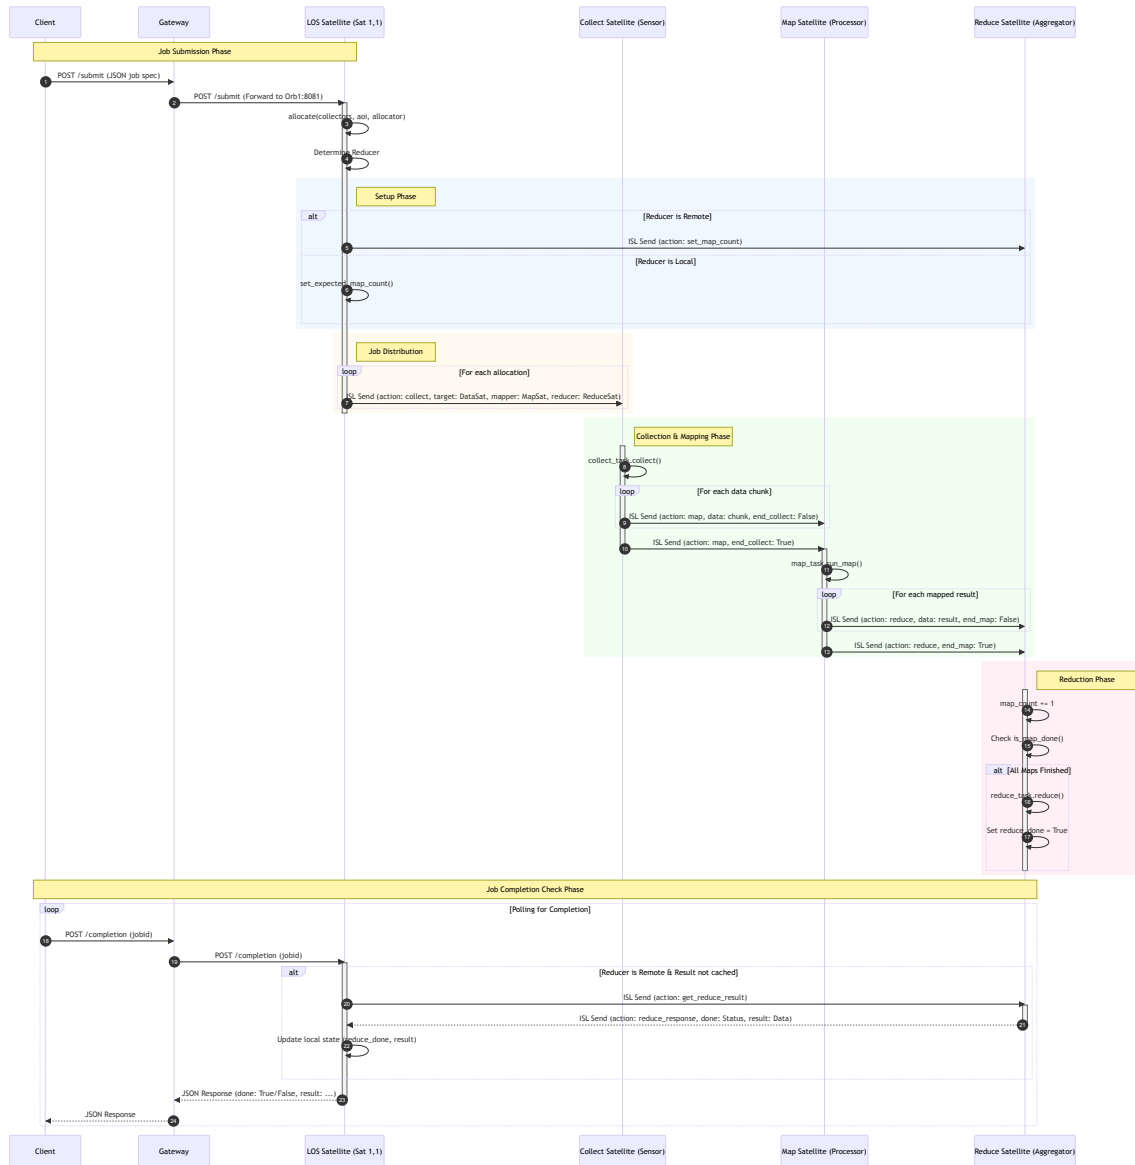
Figure 2: Job Submission and Execution Sequence

The model we use is from the Facebook ssv2 benchmark `https://github.com/facebookresearch/vjepa2`, and the default video sample used is from: `https://test-videos.co.uk/vids/bigbuckbunny/mp4/h264/360/Big_Buck_Bunny_360_10s_1MB.mp4`. More info at: `https://github.com/LeoDOS-Project/leopymr/blob/main/usecases/vjepa/README.md`

11 Docker Design

To simplify deployment, the Python simulation implementation of SpaceCoMP (LeoPyMR) places the satellite runtime in Docker containers. Custom runtimes may be deployed that extend the base Space-CoMP image with additional dependencies needed by custom Collector, Mapper and Reducer implementations.

All satellites in a constellation are deployed in their own Docker network overlay, using docker compose, where each orbital plane is represented by a service, and each satellite within that orbital plane listens on a separate port and runs a separate process within the container of the orbital plane. The rational for this design is to avoid exposing too many ports or having too many containers. Furthermore, each orbital plane may be hosted on a separate physical node using Docker Swarm `https://github.com/LeoDOS-Project/leopymr/blob/main/swarm/README.md` or Kubernetes `https://github.com/LeoDOS-Project/leopymr/blob/main/k8s/README.md`.

Requests are submitted and results are downloaded through a gateway service, which is the only service that needs to be exposed outside the cluster to the host machine. All satellites are processing requests sequentially in a single thread but all communication (ISL) is asynchronous, which mimics other platforms such as NASA cFS.

12 More Info

For more information and to cite this work please consult the paper at: `https://arxiv.org/abs/2601.17589` and use the citation:

```
@article{sandholm2026lightspeed,
  title={Lightspeed Data Compute for the Space Era},
  author={Sandholm, Thomas and Huberman, Bernardo A and Segeljakt, Klas and \
  Carbone, Paris},
  journal={arXiv preprint arXiv:2601.17589},
  year={2026}
}
```